

Introduction to Large Language Models

Week-4 Assignment

Number of questions: 10

Total mark: 10 X 1 = 10

Question 1:

A one-hot vector representation captures semantic similarity between related words like "king" and "queen".

- a) True
- b) False

Answer: b

Solution: One-hot vectors are orthogonal; no similarity is encoded.

Question 2:

Which method is used to reduce the dimensionality of a term-context matrix in count-based word representations?

- a) Principal Component Analysis
- b) Matrix Inversion
- c) Singular Value Decomposition (SVD)
- d) Latent Dirichlet Allocation

Answer: c

Solution: SVD is used to obtain low-dimensional representations in latent semantic analysis.

Question 3:

Which property makes tf-idf a better representation than raw term frequency?

- a) It is non-linear
- b) It accounts for the informativeness of words
- c) It penalizes longer documents
- d) It uses hierarchical clustering

Answer: c

Solution: IDF downweights common terms like "the" and emphasizes rare but important ones.

Question 4:

What is the purpose of using negative sampling in Word2Vec training?

- a) To reduce dimensionality of word vectors
- b) To ensure gradient convergence
- c) To balance class distribution in classification
- d) To simplify softmax computation

Answer: d

Solution: Negative sampling avoids computing softmax over the entire vocabulary.

Question 5:

In skip-gram Word2Vec, the model:

- a) Predicts a word given its context
- b) Predicts the next sentence
- c) Predicts surrounding context words given a target word
- d) Learns n-gram frequencies

Answer: c

Solution: Skip-gram learns by predicting surrounding words given a center word.

Question 6:

Why does SVD-based word embedding struggle with adding new words to the vocabulary?

- a) It uses online learning
- b) It lacks semantic interpretability
- c) It assumes word order
- d) It is computationally expensive to retrain

Answer: d

Solution: New words require recomputing the entire decomposition.

Question 7:

Which of the following best describes the term “distributional hypothesis” in NLP?

- a) Words with high frequency have greater meaning
- b) Words are defined by their part-of-speech tags
- c) A word’s meaning is characterized by the words around it
- d) Words should be normalized before vectorization

Answer: c

Question 8:

In Word2Vec, similarity between word vectors is computed using Euclidean distance.

- a) True
- b) False

Answer: b

Solution: Similarity is computed using dot product or cosine similarity.

Question 9:

Which method solves the problem of OOV (Out-Of-Vocabulary) words better?

- a) One-hot encoding
- b) CBOW
- c) Skip-gram with subsampling
- d) FastText embedding

Answer: d

Solution: FastText builds embeddings using character n-grams and handles unseen words.

Question 10:

If the word "economy" occurs 4 times in a corpus, and "growth" appears in a window of 5 words around it 3 times, what is the entry for (economy, growth) in a term-context matrix?

- a) 1
- b) 2
- c) 3
- d) 4

Answer: c

Solution: It counts co-occurrences in the window — here, 3 times.
