

Introduction to Large Language Models

Week-6 Assignment

Number of questions: 8

Total mark: $6 \times 1 + 2 \times 2 = 10$

Question 1: [1 mark]

True or False:

RoPE uses additive embeddings like sinusoidal encoding.

Answer: False

Solution: Please refer to slides.

Question 2: [1 mark]

Which of the following is true about *multi-head attention*?

- a. It increases model interpretability by using a single set of attention weights
- b. Each head operates on different parts of the input in parallel
- c. It reduces the number of parameters in the model
- d. Heads are averaged before applying the softmax function

Answer: b

Solution: Each attention head processes different learned projections of the input, enabling the model to capture different features.

Question 3: [1 mark]

What is the role of the residual connection in the Transformer architecture?

- a. Improve gradient flow during backpropagation
- b. Normalize input embeddings
- c. Reduce computational complexity
- d. Prevent overfitting

Answer: a

Solution: Please refer to lecture slides.

Question 4: [1 mark]

True or False:

The feedforward network in a Transformer block introduces non-linearity between attention layers.

Answer: True

Solution: Please refer to lecture slides.

Question 5: [1 mark]

Fill in the blank:

The sinusoidal positional encoding uses sine for even dimensions and ___ for odd dimensions.

- a. sine
- b. cosine
- c. tangent
- d. None of these

Answer: b

Solution: Please refer to lecture slides.

Question 6: [1 mark]

Why is positional encoding added to input embeddings in Transformers?

- a. To provide unique values for each word
- b. To indicate the position of tokens since Transformers are non-sequential
- c. To scale embeddings
- d. To avoid vanishing gradients

Answer: b

Solution: Please refer to lecture slides.

Question 7: [2 marks]

You are given a self-attention layer with input dimension 512, using 8 heads. What is the output dimension per head?

- a. 64
- b. 128

- c. 32
d. 256

Answer: a

Solution: Each head processes $512/8 = 64$ dimensions

QUESTION 8: [2 marks]

For a transformer with $d_{model} = 512$, calculate the positional encoding for position $p=14$ and dimensions 6 and 7 using the sinusoidal formula:

$$PE(p, 2i) = \sin\left(\frac{p}{10000^{2i/d_{model}}}\right) \quad PE(p, 2i + 1) = \cos\left(\frac{p}{10000^{2i/d_{model}}}\right)$$

- a. $\sin\left(\frac{14}{10000^{3/256}}\right), \cos\left(\frac{14}{10000^{3/256}}\right)$
- b. $\cos\left(\frac{14}{10000^{6/256}}\right), \sin\left(\frac{14}{10000^{7/256}}\right)$
- c. $\cos\left(\frac{14}{10000^{3/256}}\right), \sin\left(\frac{14}{10000^{3/256}}\right)$
- d. $\sin\left(\frac{14}{10000^{3/512}}\right), \cos\left(\frac{14}{10000^{3/256}}\right)$

Correct Answer: a

Solution:

$$\text{For dimension 6, } PE(14,6) = \sin\left(\frac{14}{10000^{6/512}}\right) = \sin\left(\frac{14}{10000^{3/256}}\right)$$

$$\text{For dimension 7, } PE(14,7) = \cos\left(\frac{14}{10000^{6/512}}\right) = \cos\left(\frac{14}{10000^{3/256}}\right)$$
